

Ridge Regression

December 1, 2018

```
In [1]: %matplotlib inline
```

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

def main():
    x1 = np.random.uniform(low=-2, high=10, size=12)
    x_train = get_x_train(x1)
    y_train = get_y_train(x1)

    print_training_data(x1, y_train)

    w = linear_regression(x_train, y_train)
    best_r = determine_reg_param(x_train, y_train)
    w_reg = ridge_regression(x_train, y_train, regularizer=best_r)

    print_regression_eq('without reg', w, round(calculate_error(w, x_train, y_train), 2))
    print()
    print_regression_eq('with reg', w_reg, round(calculate_error(w_reg, x_train, y_train), 2))

    plot_exp(x1, y_train, w, w_reg)

def determine_reg_param(x_train, y_train):
    regularizers = np.array([0.1, 1, 10, 100])
    cv_error = np.zeros_like(regularizers)
    ein = np.zeros_like(regularizers)

    for i, r in enumerate(regularizers):
        w = ridge_regression(x_train, y_train, regularizer=r)
        ein[i] = calculate_error(w, x_train, y_train)
        cv_error[i] = cross_validation(x_train, y_train, 3, r)

    print_cross_validation_results(regularizers, ein, cv_error)

    return regularizers[np.argmin(cv_error)]
```

```

def print_cross_validation_results(regularizers, ein, cv_error):
    cv_summary = np.round(np.column_stack((regularizers, ein, cv_error)), 2)
    df = pd.DataFrame(cv_summary, columns=['lambda', 'in-sample error', 'cv error'])
    print(str(df))
    print('choose lambda = ' + str(regularizers[np.argmin(cv_error)]) + '\n')

def cross_validation(x_train, y_train, num_splits, regularizer):
    d = np.column_stack((x_train, y_train))
    cv_error = np.zeros(num_splits)

    for i, this_d in enumerate(np.split(d, num_splits)):
        x_train, y_train = this_d[:, :-1], this_d[:, -1]
        w = ridge_regression(x_train, y_train, regularizer)
        cv_error[i] = calculate_error(w, x_train, y_train)

    return np.average(cv_error)

def calculate_error(w, x_train, y_train):
    y = np.dot(x_train, w)
    return np.average(np.square(np.subtract(y, y_train)))

def ridge_regression(x_train, y_train, regularizer):
    I = np.identity(x_train.shape[1])
    xTx = np.dot(x_train.T, x_train)
    xTy = np.dot(x_train.T, y_train)
    return np.dot(np.linalg.pinv(np.add(xTx, regularizer * I)), xTy)

def linear_regression(x_train, y_train):
    return np.dot(np.linalg.pinv(x_train), y_train)

def get_x_train(x1):
    bias = np.ones_like(x1)
    return np.column_stack((bias, x1))

def get_y_train(x1):
    return np.square(x1) + 10

def print_training_data(x1, y_train):
    d = np.column_stack((x1, y_train))

```

```

d_sorted = d[np.argsort(d[:, 0])]
df = pd.DataFrame(np.round(d_sorted, 2), columns=['x', 'y'])
print(str(df) + '\n')

def print_regression_eq(label, w, error):
    print('{0:12}: y = {1}x + {2}'.format(label, round(w[1], 2), round(w[0], 2)))
    print('{0:12}: {1}'.format('error', error))

def plot_exp(x1, y_train, w, w_reg):
    plt.style.use('seaborn-whitegrid')
    fig, ax = plt.subplots()

    ax.set(title='Ridge Regression')
    ax.scatter(x1, y_train, color='b', marker='x', label='training data')
    ax.plot(x1, line(w[1], x1, w[0]), color='g', label='regression line')
    ax.plot(x1, line(w_reg[1], x1, w_reg[0]), color='r', label='ridge regression line')

    ax.legend(facecolor='w', fancybox=True, frameon=True, edgecolor='black', borderpad=5)
    plt.show()

def line(m, x, b):
    return m * x + b

if __name__ == '__main__':
    main()

```

	x	y
0	-0.60	10.36
1	0.37	10.14
2	0.49	10.24
3	0.67	10.45
4	1.42	12.00
5	1.59	12.52
6	2.33	15.45
7	4.96	34.60
8	5.42	39.40
9	5.77	43.34
10	6.71	55.01
11	7.58	67.49

	lambda	in-sample error	cv error
0	0.1	24.81	15.63
1	1.0	25.02	17.03

2	10.0	30.18	40.94
3	100.0	143.70	379.07

choose lambda = 0.1

without reg : $y = 6.95x + 5.47$
error : 24.8

with reg : $y = 6.97x + 5.39$
error : 24.81

